Anthesis

Governed AI collaboration for the software development lifecycle

Anthesis is a governance-first, Git-native system for using AI in software delivery without giving up human control, reviewability, or auditability.

As AI becomes part of planning, coding, documentation, validation, and workflow execution, teams face a new operational problem: work can happen faster than it can be meaningfully reviewed, approved, or reconstructed. That creates risk around authority, trust, compliance, and accountability.

Anthesis addresses that problem by treating agent activity as governed execution rather than informal automation. Instead of asking teams to trust opaque model behavior, it forces consequential intervention into explicit, reviewable paths tied to policy, approvals, evidence, and replay-aware records.

The problem

- Agent changes that are weakly reviewed or not meaningfully attributable
- Unclear authority over who or what approved a consequential step
- Difficulty reconstructing how an outcome was produced
- Poor auditability in regulated or security-sensitive environments
- Drift between policy, intent, execution, and repository state

What Anthesis changes

From: “the agent made a change”

To: “the system executed a bounded action under policy, with recorded authority, linked approvals, evidence, and a replay-aware record.”

Core capabilities

- **Policy-bound execution** - actions operate within explicit rules, scopes, and approval requirements
- **Reviewable intervention** - consequential outputs can be shaped into forms that humans can inspect and approve
- **Replay-aware evidence** - executions can be reconstructed well enough to understand the decision surface, context, and material outcome
- **Auditable records** - approvals, actors, evidence, and outcomes are retained as governance artifacts
- **Human authority with bounded agent participation** - humans remain governance anchors for promotion, exceptions, and high-consequence decisions, while agents participate as bounded operators

Where it applies

- Code generation and repository modification
- Documentation and specification workflows
- CI/CD and workflow execution
- Validation, policy checks, and approval routing
- Audit, replay, and post-hoc reconstruction of meaningful interventions

Why this matters

- Safer automation by narrowing agent blast radius
- Better review by exposing intent, policy, authority, and supporting evidence
- Stronger incident analysis through replay-aware records
- More defensible AI adoption without discarding governance expectations
- A better fit for regulated and security-sensitive work

Ecosystem relevance

Anthesis is not only a product thesis. It is also an exploration of governance patterns for AI-era software delivery.

- Policy-governed intervention
- Human and agent actors with distinct authority models
- Evidence and replay requirements for consequential AI use
- Hardened control surfaces in the SDLC
- Reviewable and auditable automation

In one question

> **Was the system allowed to do this—and can it prove it?**

In one sentence

> **Anthesis enables human and AI collaboration in the SDLC through policy-bound, reviewable, replayable, and auditable execution, so teams can gain automation without losing control.**